

Machine Learning for Partial Differential Equations: Solvers and Operator Learning

Matthieu Darcy¹

¹California Institute of Technology

University of Bern - Institute of Mathematical Statistics and Actuarial Science
January, 2026

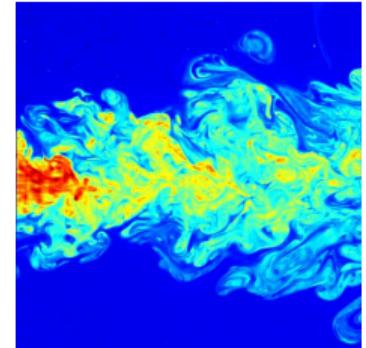
Table of Contents

- 1 Introduction: PDEs and machine learning
- 2 Operator Learning for PDEs
 - Gaussian processes/kernels for operator learning
 - Operator learning at machine precision

Introduction: PDEs and Machine learning

Many physical phenomena are modeled by PDEs:

- Fluid mechanics
- Weather prediction
- Geophysics
- ...

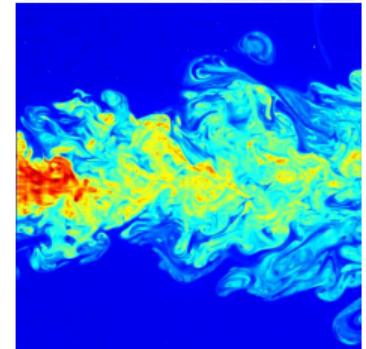


Introduction: PDEs and Machine learning

Many physical phenomena are modeled by PDEs:

- Fluid mechanics
- Weather prediction
- Geophysics
- ...

**What are the current applications of machine learning
in the context of PDEs?**



Applications of machine learning to differential equations

What are the current applications of machine learning in the context of PDEs?

- **Solve differential equations:** provide a solution to the PDE.
- **Learn differential equations:** recover unknowns of the system or learn the functional form.
- **Accelerate computations:** replace or improve existing solvers.

Introduction: partial differential equations

Partial differential equations model systems and encode laws:

$$\underbrace{\frac{\partial v}{\partial t}}_{\text{Time evolution}} = \underbrace{\nabla \cdot (a \nabla v)}_{\text{Diffusion}} + \underbrace{R(v)}_{\text{Reaction}} + \underbrace{f}_{\text{External forcing}} + \dots$$

- v is the “quantity of interest” → **solution to the PDE**
- a is a property of the medium → **coefficient**
- ...

Introduction: partial differential equations

Partial differential equations model systems and encode laws:

$$\underbrace{\frac{\partial v}{\partial t}}_{\text{Time evolution}} = \underbrace{\nabla \cdot (a \nabla v)}_{\text{Diffusion}} + \underbrace{R(v)}_{\text{Reaction}} + \underbrace{f}_{\text{External forcing}} + \dots$$

- v is the “quantity of interest” → **solution to the PDE**
- a is a property of the medium → **coefficient**
- ...

Solving the PDE gives the solution v which describes the system, but...

- How do we solve the PDE?
- What if a, f, \dots are unknown?
- What if solving the system is too expensive?
- What if we don’t know the PDE to begin with?

Introduction: partial differential equations

Partial differential equations model systems and encode known laws. Abstractly:

$$\mathcal{F}(v; a)(x) = f(x) \quad x \in \Omega$$

$$\mathcal{B}(v; a)(x) = g(x) \quad x \in \partial\Omega$$

Introduction: partial differential equations

Partial differential equations model systems and encode known laws. Abstractly:

$$\mathcal{F}(v; a)(x) = f(x) \quad x \in \Omega$$

$$\mathcal{B}(v; a)(x) = g(x) \quad x \in \partial\Omega$$

Terminology:

\mathcal{F} differential operator
 \mathcal{B} boundary operator

Unified as a single operator \mathcal{F} (laws, conserved quantities...)

v solution (want to find)

a coefficient, f forcing, g boundary/initial condition... (data, properties...)

geometry?

...

Two classes of PDE problems

Forward problem: given $\mathcal{F}, a, f, g \dots$ find v

Inverse problem: given \mathcal{F} (+ partial information) find $v, a \dots$

Learning differential equations: learn \mathcal{F} from observations

Solving and learning DEs

Machine learning solvers

- Gaussian processes [Chen, Hosseini, Owhadi, and Stuart 2021]
- Neural networks [Raissi, Perdikaris, and G. Karniadakis 2019]

parametrize a **single** function: solution v , coefficient $a \dots$

“Learn and predict a single system.”

Two classes of PDE problems

Given an operator \mathcal{S} associated to the PDE
and snapshots $(u_i, \mathcal{S}(u_i))_{i=1}^N \rightarrow \text{learn } \mathcal{S}$

} Operator learning

Neural networks

- Fourier Neural operators [Li, Kovachki, Azizzadenesheli, Liu, Bhattacharya, Stuart, and Anandkumar 2020]
- DeepONet [Lu, Jin, Pang, Zhang, and G. E. Karniadakis 2021]

parametrize an **operator** that can produce **many** solutions.

“Learn and predict across systems.”

Example: Darcy flow

Example

$$\begin{aligned}-\nabla \cdot (a(x)\nabla v(x)) &= f(x) \\ v(x) &= g(x)\end{aligned}$$

Forward problem: given a single $a, f, g \dots$ recover v .

Inverse problem: given f, g and data on v , recover a, v

Example: Darcy flow

Example

$$\begin{aligned}-\nabla \cdot (a(x)\nabla v(x)) &= f(x) \\ v(x) &= g(x)\end{aligned}$$

Forward problem: given a single $a, f, g \dots$ recover v .

Inverse problem: given f, g and data on v , recover a, v

In this talk:

Operator learning problem: learn the solution operator $\mathcal{S} : (a, f, g \dots) \mapsto v$.

Table of Contents

- 1 Introduction: PDEs and machine learning
- 2 Operator Learning for PDEs
 - Gaussian processes/kernels for operator learning
 - Operator learning at machine precision

Outline

- 1 Introduction: PDEs and machine learning
- 2 Operator Learning for PDEs
 - Gaussian processes/kernels for operator learning
 - Operator learning at machine precision

Gaussian processes/kernels for operator learning

Relevant publication: Batlle, D, Hosseini, and Owhadi “Kernel methods are competitive for operator learning”, *Journal of Computational Physics*, 2024

The operator learning problem

The operator learning problem (abstract version)

Let $\{u_i, v_i\}_{i=1}^N$ be N elements of $\mathcal{U} \times \mathcal{V}$ such that

$$\mathcal{S}(u_i) = v_i, \quad \text{for } i = 1, \dots, N.$$

The data driven operator learning problem is summarized as :

Given the data $\{u_i, v_i\}_{i=1}^N$ approximate \mathcal{S} .

\mathcal{U} is a space of functions $u : \Omega \rightarrow \mathbb{R}$

\mathcal{V} is a space of functions $v : D \rightarrow \mathbb{R}$

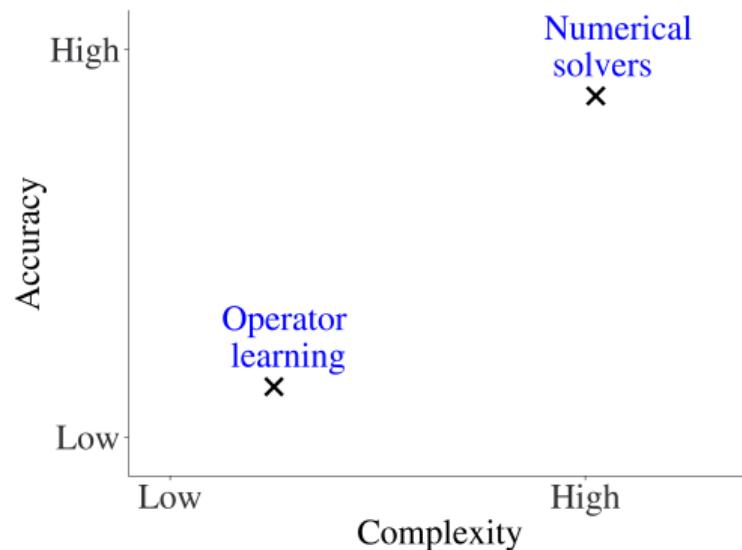
The functions u, v are *continuous, infinite dimensional* objects.

Applications of data-driven operator learning.

- **Problem 1:** unknown operator (unknown/incomplete physics, no model \mathcal{F}): learn from data. *Accuracy* is the relevant metric.

Applications of data-driven operator learning.

- **Problem 1:** unknown operator (unknown/incomplete physics, no model \mathcal{F}): learn from data. *Accuracy* is the relevant metric.
- **Problem 2** available but expensive model: accelerate computations.
Complexity-accuracy is the relevant metric.



Operator learning

Past work has focused on Operator Neural Networks that generalize Neural Networks to functional inputs and outputs:

- FNO [Li, Kovachki, Azizzadenesheli, Liu, Bhattacharya, Stuart, and Anandkumar 2020]
- PCA-net[Bhattacharya, Hosseini, Kovachki, and Stuart 2021]
- DeepOnet [Lu, Jin, Pang, Zhang, and G. E. Karniadakis 2021]
- ...

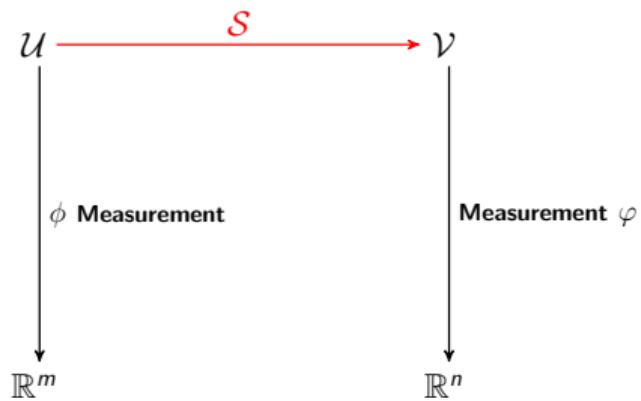
Our contribution

We propose a family of kernel/Gaussian process based-methods that are **simple, fast** and **competitive in accuracy**.

General framework

In practice we only have access to finite dimensional measurements of (u, v) given by $(\phi(u), \varphi(v))$

Encoder ϕ, φ



General framework

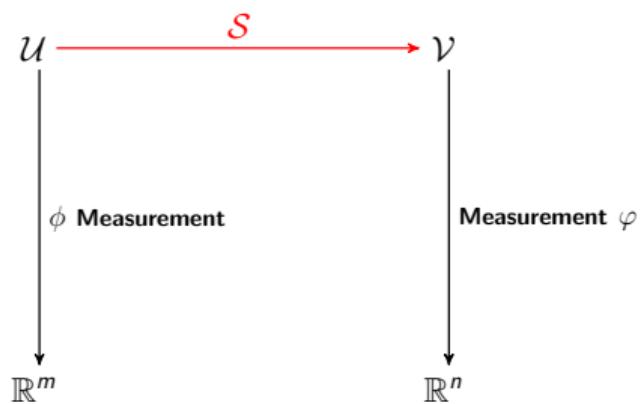
In practice we only have access to finite dimensional measurements of (u, v) given by $(\phi(u), \varphi(v))$

Encoder ϕ, φ

Common examples are:

$\phi : u \mapsto (u(x_1), \dots, u(x_m))$ Pointwise values.

$\phi : u \mapsto (\langle u, e_1 \rangle, \dots, \langle u, e_m \rangle)$ Projection in a basis
(PCA/POD, Fourier...)



General framework

In practice we only have access to finite dimensional measurements of (u, v) given by $(\phi(u), \varphi(v))$

Encoder ϕ, φ

Common examples are:

$\phi : u \mapsto (u(x_1), \dots, u(x_m))$ Pointwise values.

$\phi : u \mapsto (\langle u, e_1 \rangle, \dots, \langle u, e_m \rangle)$ Projection in a basis
(PCA/POD, Fourier...)

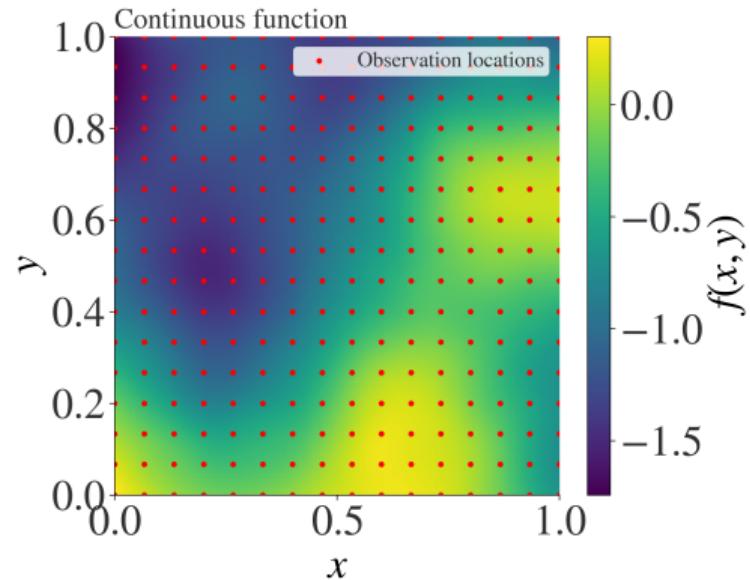


Figure: Continuum and observations

General framework

We have access to finite dimensional measurements of (u, v) given by $(\phi(u), \varphi(v))$:

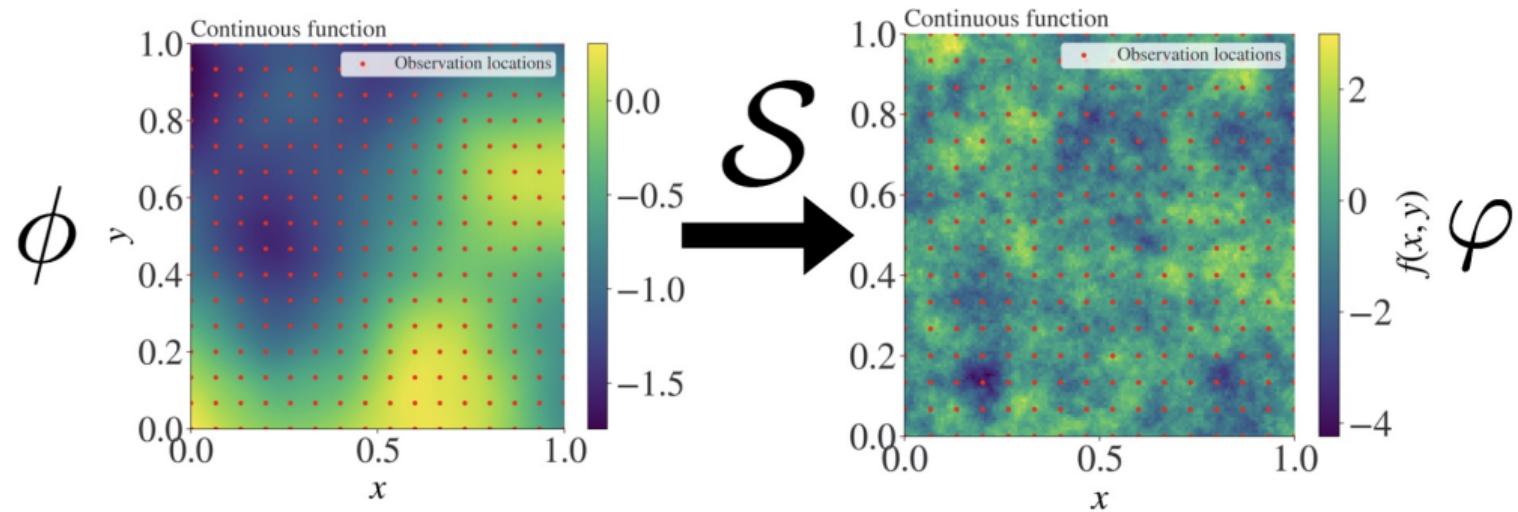


Figure: Continuum and observations

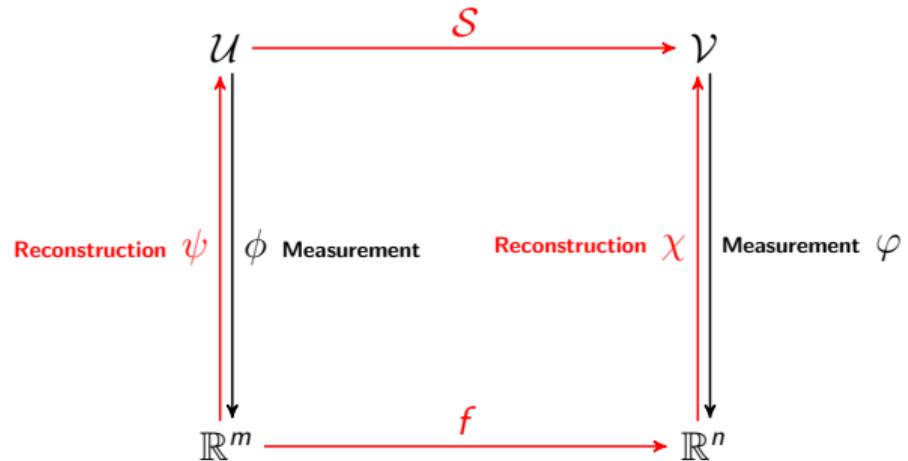
General framework

Encoder

ϕ, φ

Decoder

ψ, χ



General framework

Encoder

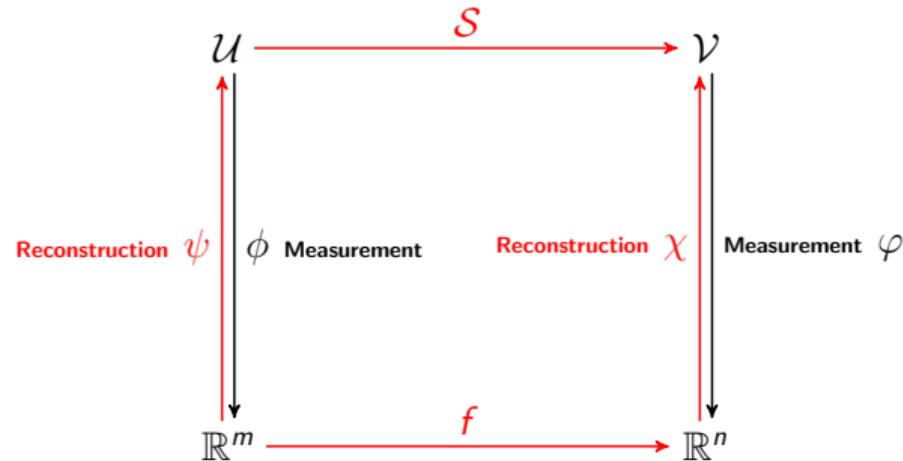
$$\phi, \varphi$$

Decoder

$$\psi, \chi$$

Finite dimensional map

$$f := \varphi \circ \mathcal{S} \circ \psi$$



General framework

Encoder

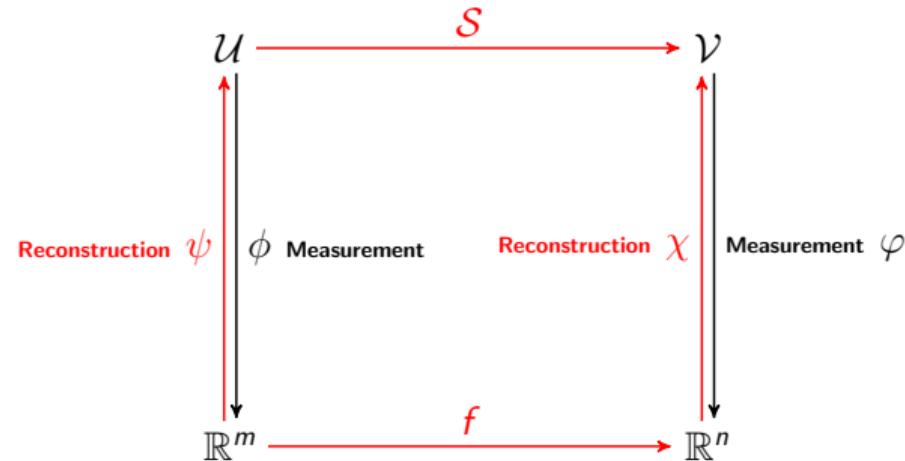
$$\phi, \varphi$$

Decoder

$$\psi, \chi$$

Finite dimensional map

$$f := \varphi \circ \mathcal{S} \circ \psi$$



Summary of our method

- ① Define the reconstructions ψ and χ as optimal recovery maps.
- ② Approximate the function f using a kernel method/GP (coordinate wise).

Step 1: Optimal recovery

We assume that \mathcal{U} and \mathcal{V} are RKHSs arising from kernels Q and K respectively. The reconstruction operators are defined as optimal recovery maps

$$\psi(\phi(u)) := \arg \min_{w \in \mathcal{U}} \|w\|_Q \quad \text{s.t.} \quad \phi(w) = \phi(u),$$

$$\chi(\varphi(v)) := \arg \min_{w \in \mathcal{V}} \|w\|_K \quad \text{s.t.} \quad \varphi(w) = \varphi(v),$$

The maps are the minmax optimal recovery of u and v respectively [Owhadi and Scovel].

Step 1: Optimal recovery

We assume that \mathcal{U} and \mathcal{V} are RKHSs arising from kernels Q and K respectively. The reconstruction operators are defined as optimal recovery maps

$$\psi(\phi(u)) := \arg \min_{w \in \mathcal{U}} \|w\|_Q \quad \text{s.t.} \quad \phi(w) = \phi(u),$$

$$\chi(\varphi(v)) := \arg \min_{w \in \mathcal{V}} \|w\|_K \quad \text{s.t.} \quad \varphi(w) = \varphi(v),$$

The maps are the minmax optimal recovery of u and v respectively [Owhadi and Scovel].

Optimal recovery maps can be expressed by solving a linear system

$$\psi(\phi(u))(x) = Q(x, X)Q(X, X)^{-1}\phi(u) \quad \text{and} \quad \chi(\varphi(v))(y) = K(y, Y)K(Y, Y)^{-1}\varphi(v).$$

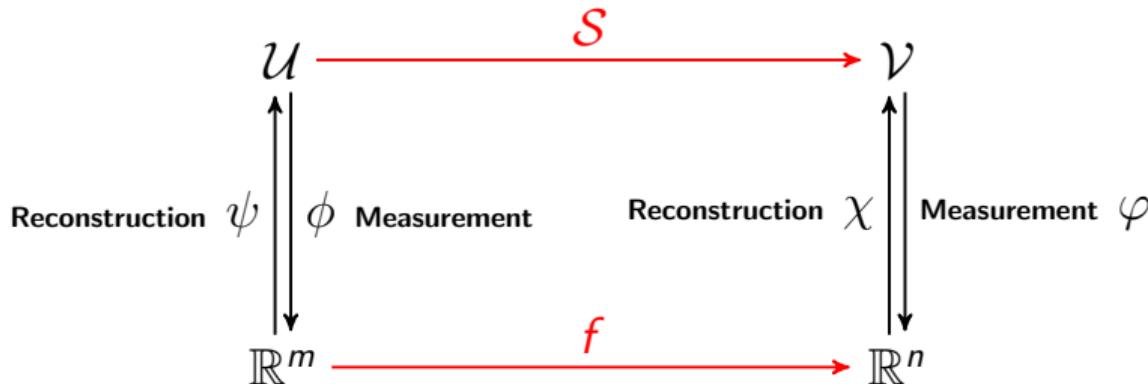
Step 2: Recovery of f

Once the reconstruction operators ψ and χ are defined, our best strategy is to reconstruct f in the diagram:

$$f^\dagger \approx f := \varphi \circ \mathcal{S} \circ \psi$$

and to approximate the operator \mathcal{S} with the operator

$$\mathcal{S}^\dagger := \chi \circ f^\dagger \circ \phi.$$



A simple Gaussian process method for f^\dagger

Given a kernel S , we approximate $f : \mathbb{R}^m \rightarrow \mathbb{R}^n$ via optimal recovery **independently component wise**:

$$f_j^\dagger := \arg \min_{h \in \mathcal{H}_S} \|h\|_S \quad \text{s.t.} \quad h(\phi(u_i)) = (\varphi(v_i))_j \quad \text{for } i = 1, \dots, N.$$

A simple Gaussian process method for f^\dagger

Given a kernel S , we approximate $f : \mathbb{R}^m \rightarrow \mathbb{R}^n$ via optimal recovery **independently component wise**:

$$f_j^\dagger := \arg \min_{h \in \mathcal{H}_S} \|h\|_S \quad \text{s.t.} \quad h(\phi(u_i)) = (\varphi(v_i))_j \quad \text{for } i = 1, \dots, N.$$

This also has a closed-form solution given a linear solve:

$$f_j^\dagger(\mathbf{u}) = S(\mathbf{u}, U)S(U, U)^{-1}\mathbf{v}_j.$$

where $U_i := \phi(u_i)$ and $V_i := \varphi(v_i)$.

Can be solved in parallel (or even better).

Measurement invariance

Mesh invariance: translate between meshes.

$$h := \tilde{\varphi} \circ \chi \circ f \circ \phi \circ \tilde{\psi} \equiv \tilde{\varphi} \circ \mathcal{S} \circ \tilde{\psi}.$$

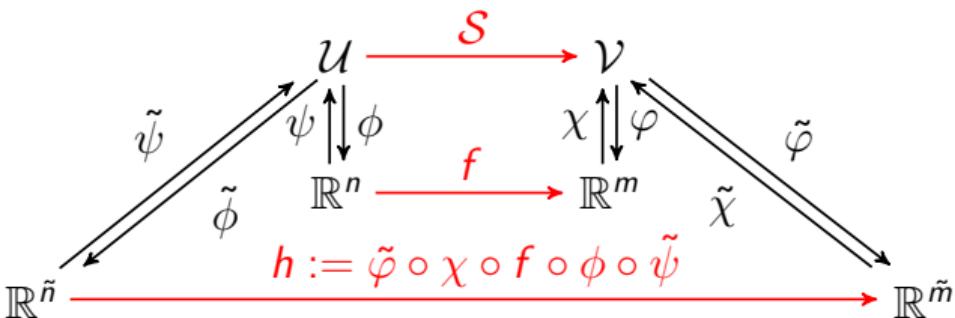


Figure: Mesh invariance of the method.

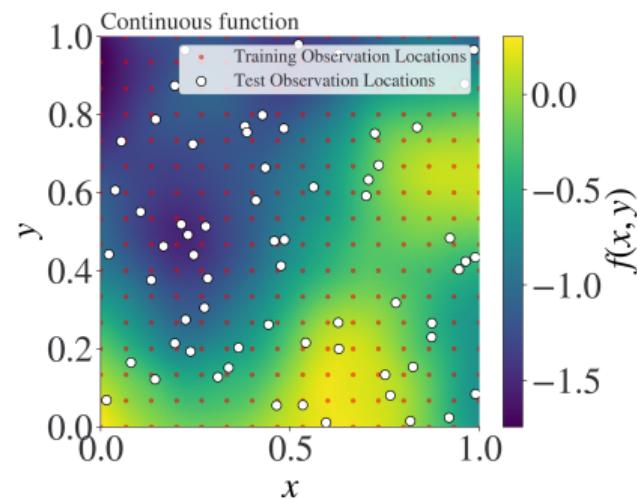


Figure: Change of mesh

Summary of the method

Summary of our method

- ① Define the reconstructions ψ and χ as the optimal recovery map (linear solve).
- ② Approximate the function f using a GP/kernel method.

Key takeways

- **Simplicity:** simplifying assumptions (independence of steps 1 and 2, coordinate wise regression of f ...) for improved complexity.
- **UQ:** Gaussian process interpretation provides UQ.
- **Theory:** convergence guarantees (new for operator learning).
- **Experimentally:** competitive in accuracy and cost/accuracy (low-medium data regime).

Summary of results: accuracy

	Low-data regime			High-data regime			
	Burger's	Darcy problem	Advection I	Advection II	Hemholtz	Structural Mechanics	Navier Stokes
DeepONet	2.15%	2.91%	0.66%	15.24%	5.88%	5.20%	3.63%
POD-DeepONet	1.94%	2.32%	0.04%	n/a	n/a	n/a	n/a
FNO	1.93%	2.41%	0.22%	13.49%	1.86%	4.76%	0.26%
PCA-Net	n/a	n/a	n/a	12.53%	2.13%	4.67%	2.65%
PARA-Net	n/a	n/a	n/a	16.64%	12.54%	4.55%	4.09%
Linear	36.24%	6.74%	$2.15 \times 10^{-13}\%$	11.28%	10.59%	27.11%	5.41%
Best of Matérn/RQ	2.15%	2.75%	$2.75 \times 10^{-3}\%$	11.44%	1.00%	5.18%	0.12%

Table: Summary of numerical results: L^2 relative test error.

We compare the performance against several neural operator frameworks on a range of problems taken from two benchmarking studies [Hoop, Huang, Qian, and Stuart 2022, Lu, Meng, Cai, Mao, Goswami, Zhang, and G. E. Karniadakis 2022] and find competitive performance.

Experimental results: Navier-Stokes

Vorticity-stream formulation of the incompressible Navier-Stokes equations:

$$\frac{\partial v}{\partial t} + (w \cdot \nabla)v - \nu \Delta v = f$$

$$v = -\Delta \psi$$

$$\int_D \psi = 0$$

$$w = \left(\frac{\partial \psi}{\partial x_2}, -\frac{\partial \psi}{\partial x_1} \right)$$

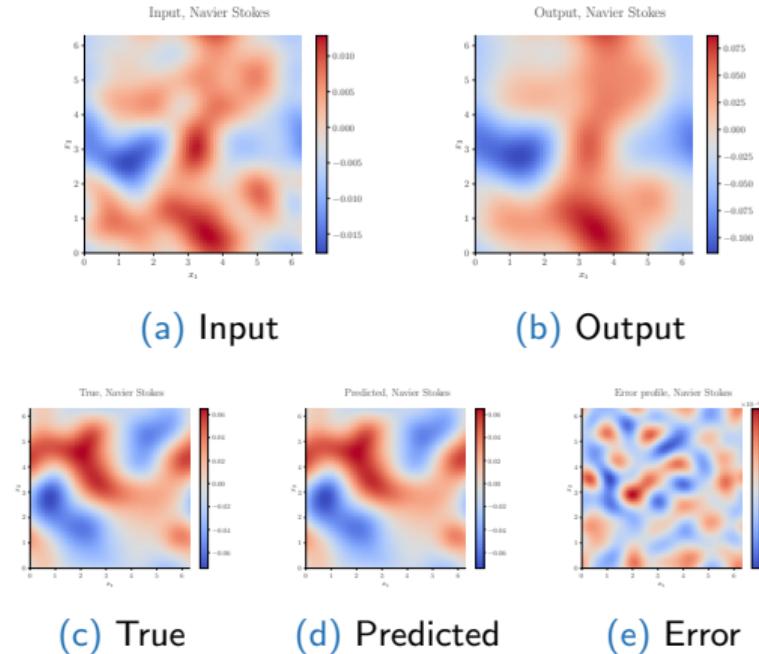
Learn the map from the forcing term f to the vorticity field v at $t = T$:

$$\mathcal{S} : f \mapsto v(\cdot, T)$$

Navier-Stokes

Method	Accuracy
DeepONet	3.63 %
FNO	0.26 %
PCA-Net	2.32 %
Linear method	5.41 %
Kernel method	0.12%

Table: L^2 relative error
on Navier-Stokes.



Empirical results

Competitive in complexity-accuracy:

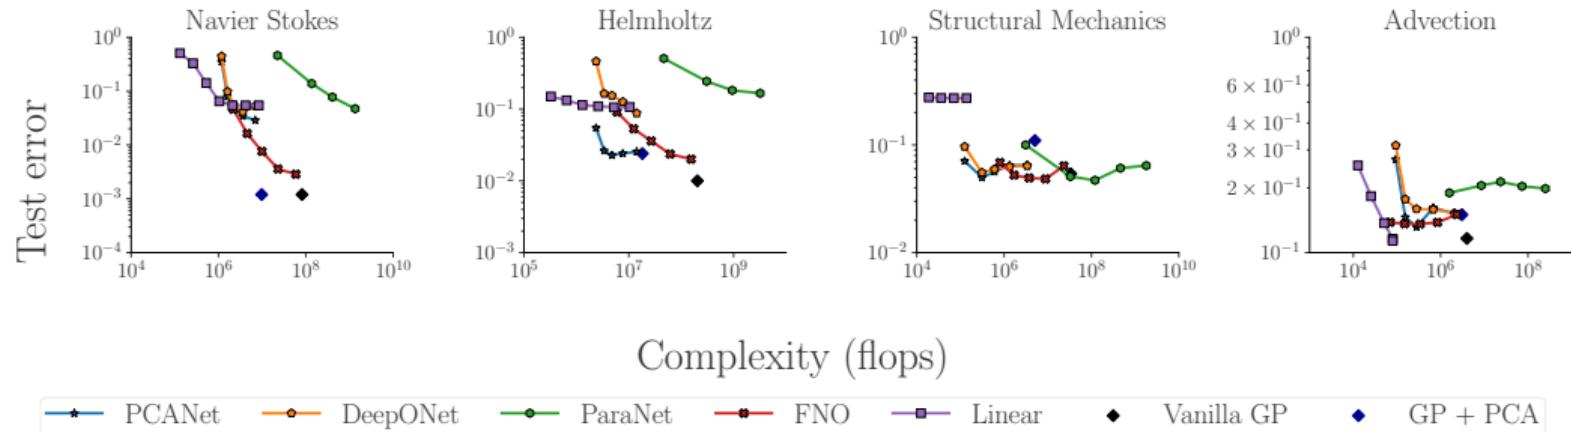


Figure: Cost-accuracy tradeoff at prediction.

Outline

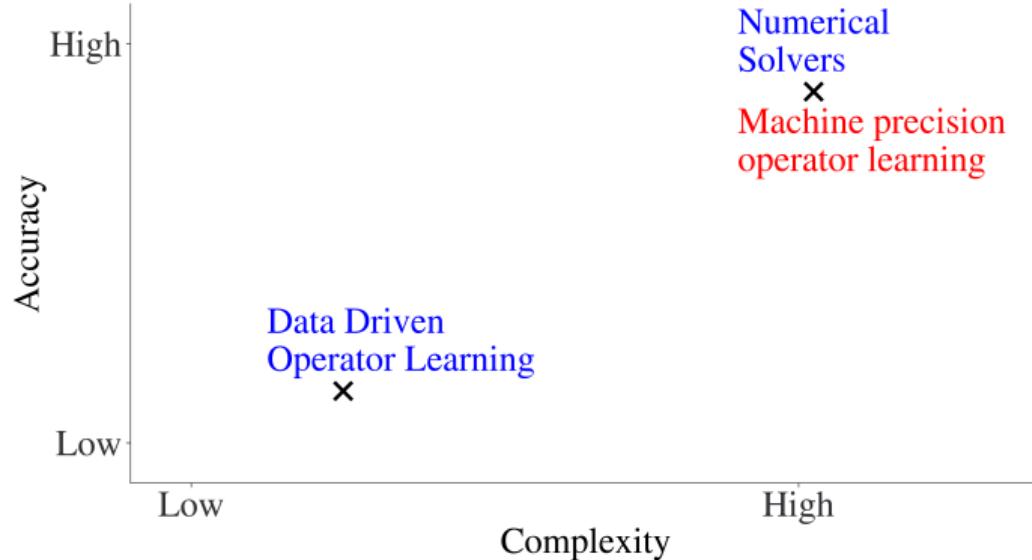
- 1 Introduction: PDEs and machine learning
- 2 Operator Learning for PDEs
 - Gaussian processes/kernels for operator learning
 - Operator learning at machine precision

High accuracy operator learning

- The presented framework is well suited for **Problem 1**: unintrusive, doesn't use any information about \mathcal{S} .

High accuracy operator learning

- The presented framework is well suited for **Problem 1**: unintrusive, doesn't use any information about \mathcal{S} .
- **Problem 2**, can operator learning match the accuracy of traditional solvers?



Operator learning at machine precision

If we know the underlying PDE operator \mathcal{F} and have access to an underlying solver, can we achieve the same high/machine precision?

Operator learning at machine precision

If we know the underlying PDE operator \mathcal{F} and have access to an underlying solver, can we achieve the same high/machine precision?

We propose CHONKORIS:a data-driven operator learning framework, that learns to emulate a (quasi)-Newton solvers and can achieve machine precision.

Relevant publication: Bacho, Sorokin, Yang, Bourdais, Calvello, D, Hsu, Hosseini, and Owhadi *Operator Learning at Machine Precision*, 2025

Operator learning at machine precision

Observation: the PDE defines a constraint for \mathcal{S} :

$$\mathcal{F}(\mathcal{S}(u), u) = 0$$

Operator learning at machine precision

Observation: the PDE defines a constraint for \mathcal{S} :

$$\mathcal{F}(\mathcal{S}(u), u) = 0$$

Ex:

$$\mathcal{F}(\mathcal{S}(u), u) = -\nabla \cdot (u \nabla v) + \tau(v) = 0$$

$$\mathcal{S}(u) = v \quad (\text{Coefficient to solution map})$$

Operator learning at machine precision

Observation: the PDE defines a constraint for \mathcal{S} :

$$\mathcal{F}(\mathcal{S}(u), u) = 0$$

Ex:

$$\begin{aligned}\mathcal{F}(\mathcal{S}(u), u) &= -\nabla \cdot (u \nabla v) + \tau(v) = 0 \\ \mathcal{S}(u) &= v \quad (\text{Coefficient to solution map})\end{aligned}$$

Observation: We can solve this using iterative methods (Newton-Raphson)

$$v_{n+1} = v_n - \left(\frac{\delta}{\delta v} \mathcal{F}(v_n, u) \right)^{-1} \mathcal{F}(v_n, u)$$

Key idea: it is sufficient to learn $\mathcal{S}(v, u) = \left(\frac{\delta}{\delta v} \mathcal{F}(v, u) \right)^{-1}$

Operator learning at machine precision

High level idea: the optimization algorithm defines an operator \mathcal{S}

$$u \text{ (PDE data)} \xrightarrow{\text{Newton flow } \mathcal{S}} v \text{ (output of the algorithm)}$$

Traditional operator learning learns the entire flow \mathcal{S}

Operator learning at machine precision

High level idea: the optimization algorithm defines an operator \mathcal{S}

$$u \text{ (PDE data)} \xrightarrow{\text{Newton flow } \mathcal{S}} v \quad (\text{output of the algorithm})$$

Traditional operator learning learns the entire flow \mathcal{S}

$$(v_0, u) \xrightarrow{\text{Iterate } \mathcal{R}} (v_1, u) \xrightarrow{\text{Iterate } \mathcal{R}} \dots (v_n, u)$$

CHONKNORIS learns the iterates of the flow \mathcal{R} .

Key takeways:

- Achieves machine precision (solver accuracy) on many problems
- In practice use Gauss-Newton (regularized Newton).
- Can be combined with previous operator learning approaches for the initial guess.
- Enables a foundation model that can generalize to unseen PDEs.

Key takeways:

- Achieves machine precision (solver accuracy) on many problems
- In practice use Gauss-Newton (regularized Newton).
- Can be combined with previous operator learning approaches for the initial guess.
- Enables a foundation model that can generalize to unseen PDEs.

Limitations

- Intrusive: requires a solver and its intermediate states.
- Still slower than the original solver.
- Limited to low dimensional examples (forming and inverting the Hessian is bad in high dimensions).
- Inverse problems are harder due to ill-conditioning.

Example

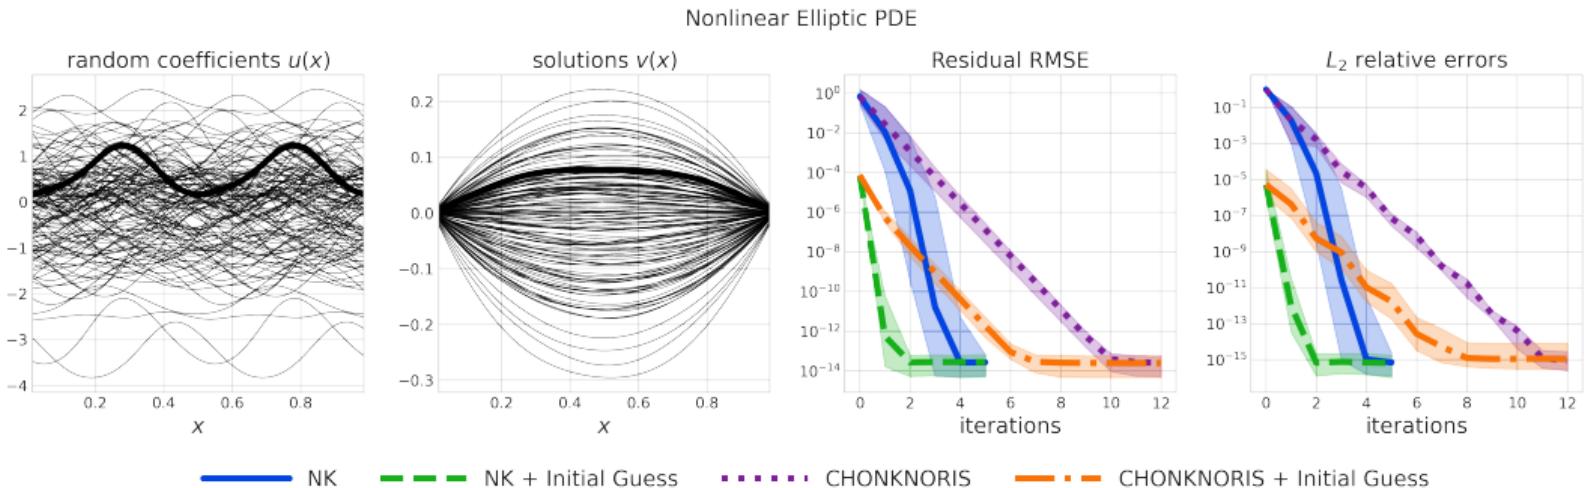


Figure: Nonlinear elliptic PDE: $10^{-5} \rightarrow 10^{-15}$ error

Summary and outlook

We proposed two methods for operator learning:

- A general purpose Gaussian process approach for inintrusive, low-accuracy data-driven operator learning.
- An intrusive operator learning framework for high-accuracy operator learning.

Summary and outlook

We proposed two methods for operator learning:

- A general purpose Gaussian process approach for inintrusive, low-accuracy data-driven operator learning.
- An intrusive operator learning framework for high-accuracy operator learning.

Outlook and future work:

- Is it possiblle to adress operator learning between **Problem 1** and **Problem 2**, i.e. partial knowledge of the physics (conserved quantities, mathematical properties...).
- How can we include geometry?
- Foundation models for PDEs: how doe we enable cross-pde learning at scale?
- Can we do this while preserving convergence guarantees?

Questions? Comments?



(a) Operator learning
with Gaussian processes



(b) Machine precision
operator learning

References I

-  Bacho, Aras et al. [2025]. *Operator Learning at Machine Precision*. arXiv: 2511.19980 [cs.LG]. URL: <https://arxiv.org/abs/2511.19980>.
-  Baptista, Ricardo et al. [2025]. *Solving Roughly Forced Nonlinear PDEs via Misspecified Kernel Methods and Neural Networks*. arXiv: 2501.17110 [math.NA]. URL: <https://arxiv.org/abs/2501.17110>.
-  Batlle, Pau et al. [2024]. “Kernel methods are competitive for operator learning”. In: *Journal of Computational Physics* 496.
-  Bhattacharya, Kaushik et al. [2021]. “Model Reduction And Neural Networks For Parametric PDEs”. en. In: *The SMAI Journal of computational mathematics* 7, pp. 121–157. DOI: 10.5802/smai-jcm.74. URL: <https://smai-jcm.centre-mersenne.org/articles/10.5802/smai-jcm.74/>.
-  Chen, Yifan et al. [2021]. “Solving and learning nonlinear PDEs with Gaussian processes”. In: *Journal of Computational Physics* 447, p. 110668. ISSN: 0021-9991. DOI: <https://doi.org/10.1016/j.jcp.2021.110668>. URL: <https://www.sciencedirect.com/science/article/pii/S0021999121005635>.

References II

-  D, Matthieu et al. [2023]. "One-shot learning of stochastic differential equations with data adapted kernels". In: *Physica D: Nonlinear Phenomena* 444, p. 133583. ISSN: 0167-2789. DOI: <https://doi.org/10.1016/j.physd.2022.133583>. URL: <https://www.sciencedirect.com/science/article/pii/S0167278922002871>.
-  Hoop, Maarten V. de et al. [2022]. *The Cost-Accuracy Trade-Off In Operator Learning With Neural Networks*.
-  Li, Zongyi et al. [2020]. "Fourier Neural Operator for Parametric Partial Differential Equations". In: *CoRR* abs/2010.08895. arXiv: 2010.08895. URL: <https://arxiv.org/abs/2010.08895>.
-  Lu, Lu et al. [Mar. 2021]. "Learning nonlinear operators via DeepONet based on the universal approximation theorem of operators". In: *Nature Machine Intelligence* 3.3, pp. 218–229.

References III

-  Lu, Lu et al. [2022]. "A comprehensive and fair comparison of two neural operators (with practical extensions) based on FAIR data". In: *Computer Methods in Applied Mechanics and Engineering* 393, p. 114778. ISSN: 0045-7825.
-  Owhadi, Houman and Clint Scovel [2019]. *Operator-Adapted Wavelets, Fast Solvers, and Numerical Homogenization: From a Game Theoretic Approach to Numerical Approximation and Algorithm Design*. Cambridge Monographs on Applied and Computational Mathematics. Cambridge University Press.
-  Raissi, M., P. Perdikaris, and G.E. Karniadakis [2019]. "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations". In: *Journal of Computational Physics* 378, pp. 686–707. ISSN: 0021-9991. DOI: <https://doi.org/10.1016/j.jcp.2018.10.045>. URL: <https://www.sciencedirect.com/science/article/pii/S0021999118307125>.

Table of Contents

- 3 Operator learning: Theoretical guarantees
- 4 Learning and solving differential equations
 - Learning SODEs in sparse data-regimes
 - Solving PDEs with stochasticity and high frequency

Theory: assumptions

Suppose that

\mathcal{U} is an RKHS of functions $u : \Omega \rightarrow \mathbb{R}$

\mathcal{V} is an RKHS of functions $v : D \rightarrow \mathbb{R}$.

Assumption (Reconstruction operators - informal)

- *Regularity of the domains Ω and D : compactness and Lipschitz boundary.*
- *Regularity of the kernels Q and K .*
- *Space filling property of collocation points.*

Assumption (Approximation of \mathcal{S} - informal)

- *Regularity of the operator \mathcal{S} : continuity and bounded Fréchet derivatives.*
- *Regularity of the kernels S^n .*
- *Resolution and space-filling property of the data.*

Theory: convergence result

Under a set of technical assumptions, we have the following theorem:

Theorem (Condensed version of Main Theorem)

For all $t' \in (0, t)$,

$$\lim_{n,m \rightarrow \infty} \lim_{N \rightarrow \infty} \sup_{u \in B_R(\mathcal{H}_Q)} \|\mathcal{S}(u) - \chi^m \circ f_N^{\dagger, m, n} \circ \phi^n(u)\|_{H^{t'}(D)} \rightarrow 0.$$

Table of Contents

- ③ Operator learning: Theoretical guarantees
- ④ Learning and solving differential equations
 - Learning SODEs in sparse data-regimes
 - Solving PDEs with stochasticity and high frequency

Outline

- ③ Operator learning: Theoretical guarantees
- ④ Learning and solving differential equations
 - Learning SODEs in sparse data-regimes
 - Solving PDEs with stochasticity and high frequency

Learning SODEs

Relevant publication: D, Hamzi, Livieri, Owhadi, and Tavallali “One-shot learning of stochastic differential equations with data adapted kernels”, *Physica D: Nonlinear Phenomena*, 2023

Learning SDEs

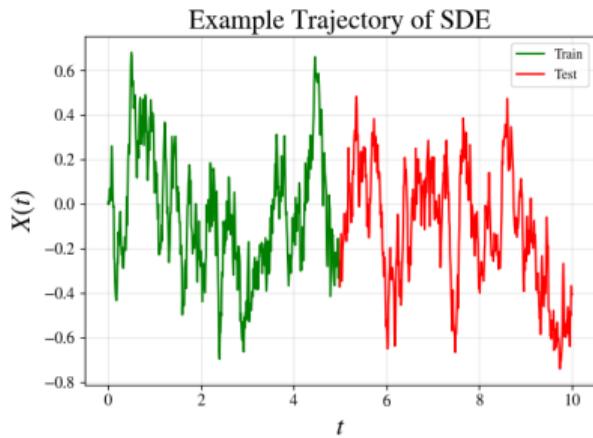
The problem

Learn the unknown drift f and the diffusion σ of a generic SODE:

$$dX_t = f(X_t) dt + \sigma(X_t) dW_t$$

from a finite number of samples $X := (X_{t_n})_{n=1}^N$.

Learn the evolution of the system from samples.



Learning SDEs

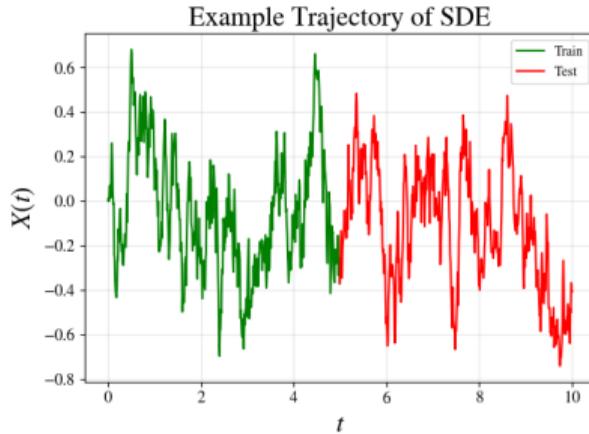
The problem

Learn the unknown drift f and the diffusion σ of a generic SODE:

$$dX_t = f(X_t) dt + \sigma(X_t) dW_t$$

from a finite number of samples $X := (X_{t_n})_{n=1}^N$.

Learn the evolution of the system from samples.



The difficulty

- The observations X come from a *single* trajectory with N small (sparse data).
- We want a flexible method: few assumptions on f and σ .
- A large sampling time-steps Δt introduces a discretization error.

Summary of the method

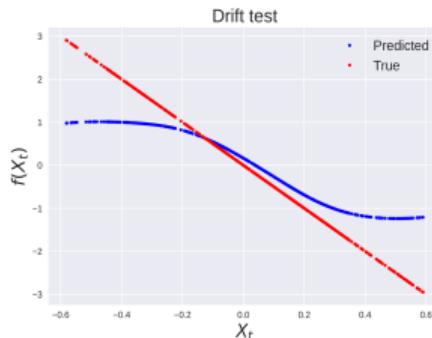
Our approach

- ① Place a Gaussian Process prior on f, σ .
- ② Formulate and solve an optimization problem using Maximum A Posteriori (MAP) Estimation.
- ③ Optimize hyper-parameters using cross-validation → enables generalization in sparse data regimes!

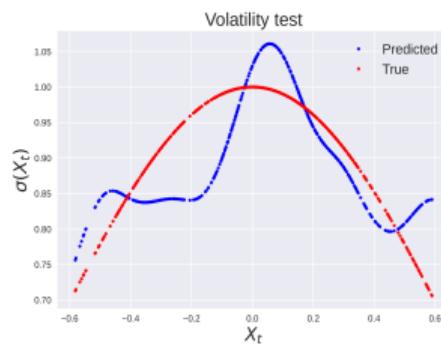
Example

Drift f

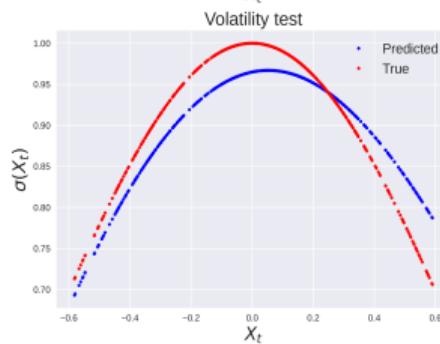
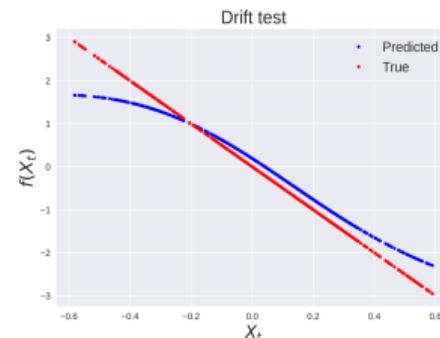
Non learned kernel



Diffusion σ



Learned kernel



Learning the right kernel makes a big difference in performance!

Outline

- ③ Operator learning: Theoretical guarantees
- ④ Learning and solving differential equations
 - Learning SODEs in sparse data-regimes
 - Solving PDEs with stochasticity and high frequency

Solving PDEs with stochasticity and high frequency

Relevant publication: Baptista, Calvello, D, Owhadi, Stuart, and Yang *Solving Roughly Forced Nonlinear PDEs via Misspecified Kernel Methods and Neural Networks*, 2025

Solving rough PDEs

The problem

We want to **develop machine learning solvers for rough PDEs** of the form

$$\begin{aligned}\mathcal{F}(v) &= \xi, & x \in \Omega, \\ v &= 0, & x \in \partial\Omega\end{aligned}\tag{RPDE}$$

Solving rough PDEs

The problem

We want to **develop machine learning solvers for rough PDEs** of the form

$$\begin{aligned}\mathcal{F}(v) &= \xi, \quad x \in \Omega, \\ v &= 0, \quad x \in \partial\Omega\end{aligned}$$

The difficulty

Solving (RPDE) is difficult because the forcing term is ‘rough’:

- ξ is stochastic/high frequency $\rightarrow \xi \in H^{-s} \neq L^2$
- v has low regularity.

Machine learning solvers struggle to capture high frequency.

Summary of the method

We modify the loss to adapt to the roughness/high frequency:

$$\arg \min_u \|\mathcal{F}(u) - \xi\|_{L^2}^2 + \|u\|_{L^2(\partial\Omega)}^2 + \gamma \mathcal{R}(u) \quad \text{Physics Informed Loss}$$

Summary of the method

We modify the loss to adapt to the roughness/high frequency:

$$\arg \min_u \|\mathcal{F}(u) - \xi\|_{L^2}^2 + \|u\|_{L^2(\partial\Omega)}^2 + \gamma \mathcal{R}(u) \quad \text{Physics Informed Loss}$$



$$\arg \min_u \|\mathcal{F}(u) - \xi\|_{H^{-s}}^2 + \|u\|_{L^2(\partial\Omega)}^2 + \gamma \mathcal{R}(u) \quad \text{Negative Sobolev Loss}$$

Summary of the method

We modify the loss to adapt to the roughness/high frequency:

$$\arg \min_u \|\mathcal{F}(u) - \xi\|_{L^2}^2 + \|u\|_{L^2(\partial\Omega)}^2 + \gamma \mathcal{R}(u) \quad \text{Physics Informed Loss}$$



$$\arg \min_u \|\mathcal{F}(u) - \xi\|_{H^{-s}}^2 + \|u\|_{L^2(\partial\Omega)}^2 + \gamma \mathcal{R}(u) \quad \text{Negative Sobolev Loss}$$



$$\arg \min_u |\mathcal{F}(u) - \xi|_{\Phi^N}^2 + \sum_{i=1}^m |u(x_i)|^2 + \gamma \mathcal{R}(u) \quad \text{Discretized loss}$$

Illustration of results

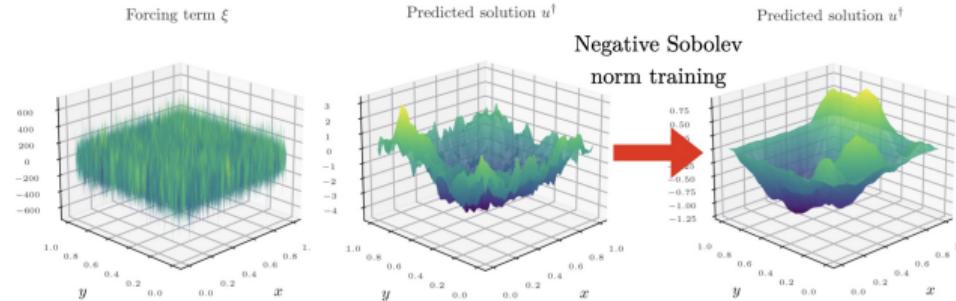


Figure: Negative Sobolev training

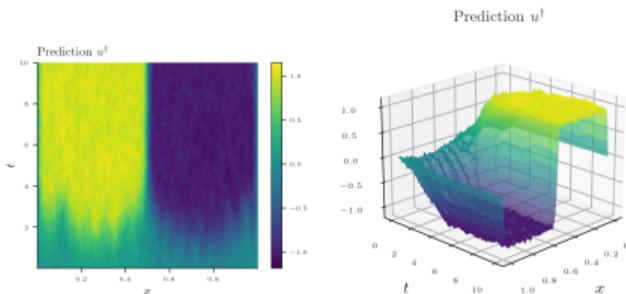


Figure: Time dependent: Stochastic Allen-Cahn equation.

Key takeways

Theoretical contributions:

- Proposes a negative Sobolev norm (related to weak solutions).
- Prove the convergence of the kernel method (even with misspecification).

Computational contributions:

- Provide an efficient numerical discretization of the Sobolev norm.
- Provide an efficient numerical method to minimize the resulting loss.
- Demonstrate the framework with neural networks and kernels.

Summary and outlook

I have worked on two learning and solving differential equations involving stochasticity:

- Learning SDEs from sparse data.
- Solving PDEs driven by rough forcing.

Summary and outlook

I have worked on two learning and solving differential equations involving stochasticity:

- Learning SDEs from sparse data.
- Solving PDEs driven by rough forcing.

Outlook and future work:

- Can we use some of these techniques for operator learning in sparse data regime
- Can we use some of these insights for machine learning in the presence of roughness/shocks/discontinuity?
- Can we provide better probabilistic predictions and more scalable approaches for GPs (work in progress).

Optimal recovery

We will assume that \mathcal{U} and \mathcal{V} are RKHSs arising from kernels Q and K respectively. The reconstruction operators are defined as optimal recovery maps

$$\psi(\phi(u)) := \arg \min_{w \in \mathcal{U}} \|w\|_Q \quad \text{s.t.} \quad \phi(w) = \phi(u),$$

$$\chi(\varphi(v)) := \arg \min_{w \in \mathcal{V}} \|w\|_K \quad \text{s.t.} \quad \varphi(w) = \varphi(v),$$

The maps are the minmax optimal recovery of u and v respectively [Owhadi and Scovel].

Optimal recovery

We will assume that \mathcal{U} and \mathcal{V} are RKHSs arising from kernels Q and K respectively. The reconstruction operators are defined as optimal recovery maps

$$\psi(\phi(u)) := \arg \min_{w \in \mathcal{U}} \|w\|_Q \quad \text{s.t.} \quad \phi(w) = \phi(u),$$

$$\chi(\varphi(v)) := \arg \min_{w \in \mathcal{V}} \|w\|_K \quad \text{s.t.} \quad \varphi(w) = \varphi(v),$$

The maps are the minmax optimal recovery of u and v respectively [Owhadi and Scovel].

Optimal recovery maps can be expressed by solving a linear system

$$\psi(\phi(u))(x) = Q(x, X)Q(X, X)^{-1}\phi(u) \quad \text{and} \quad \chi(\varphi(v))(y) = K(y, Y)K(Y, Y)^{-1}\varphi(v).$$

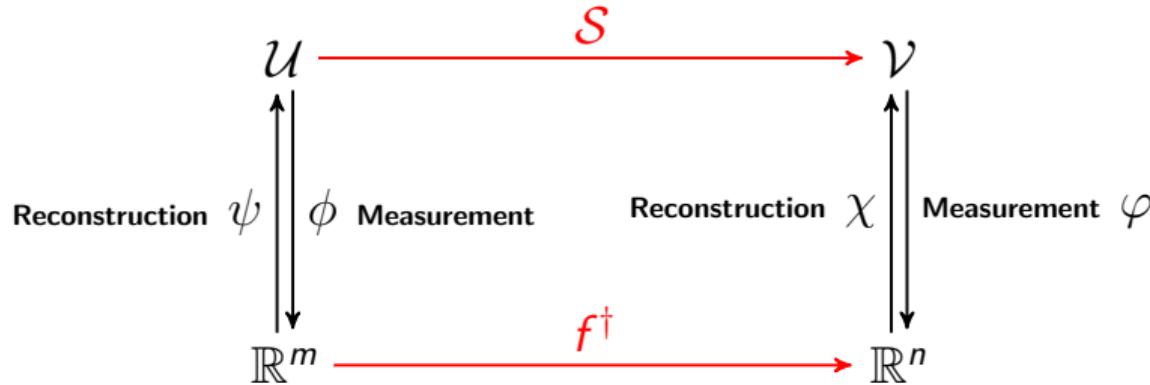
Recovery of f^\dagger

Once the reconstruction operators ψ and χ are defined, our best strategy is to reconstruct f^\dagger in the diagram:

$$\bar{f} \approx f^\dagger := \varphi \circ \mathcal{S} \circ \psi$$

and to approximate the operator \mathcal{S} with the operator

$$\bar{\mathcal{S}} := \chi \circ \bar{f} \circ \phi.$$



A simple kernel method for f^\dagger

Given a kernel S , we approximate $f : \mathbb{R}^m \rightarrow \mathbb{R}^n$ via optimal recovery **independently component wise**:

$$\bar{f}_j := \arg \min_{h \in \mathcal{H}_S} \|h\|_S \quad \text{s.t.} \quad h(\phi(u_i)) = (\varphi(v_i))_j \quad \text{for } i = 1, \dots, N.$$

A simple kernel method for f^\dagger

Given a kernel S , we approximate $f : \mathbb{R}^m \rightarrow \mathbb{R}^n$ via optimal recovery **independently component wise**:

$$\bar{f}_j := \arg \min_{h \in \mathcal{H}_S} \|h\|_S \quad \text{s.t.} \quad h(\phi(u_i)) = (\varphi(v_i))_j \quad \text{for } i = 1, \dots, N.$$

This also has a closed-form solution given a linear solve:

$$\bar{f}_j(\mathbf{u}) = S(\mathbf{u}, U)S(U, U)^{-1}\mathbf{v}_j.$$

where $U_i := \phi(u_i)$ and $V_i := \varphi(v_i)$.

Can be solved in parallel (or even better).

Advantages the simple approach

The kernel S can be a standard kernel such as the linear¹, squared exponential or Matérn kernel. This simple choice already offers several advantages:

- ① Low cost in training (< 5 seconds on a workstation) and at prediction (in the low-medium data regime).
- ② Competitive accuracy.
- ③ Empirically robust to choice of hyper-parameters/kernels.
- ④ Simple to implement: several libraries solve this problem out of the box.
- ⑤ The Gaussian process interpretation provides uncertainty quantification.
- ⑥ Convergence guarantees.

If you do operator learning, you should try using Gaussian processes!

¹Equivalent to doing linear regression

Complexity-accuracy tradeoff

We evaluate our operator learning methods through the **cost-accuracy tradeoff**:

- The **accuracy** is measured by the relative L^2 loss:

$$\mathcal{R}_N(\mathcal{S}) = \frac{1}{N} \sum_{i=1}^N \left[\frac{\|\mathcal{S}(u_i) - \mathcal{S}^\dagger(u_i)\|_{L^2}}{\|\mathcal{S}(u_i)\|_{L^2}} \right]$$

- The **complexity** depends on the **training cost** (qualitative metrics) and the **prediction cost** (measured in floating point operations - FLOPs).

We compare the test performance of our method with different choices of the kernel S using the examples from two comparison papers [Hoop, Huang, Qian, and Stuart], [Lu, Meng, Cai, Mao, Goswami, Zhang, and G. E. Karniadakis]

Navier-Stokes

Vorticity-stream formulation of the incompressible Navier-Stokes equations:

$$\frac{\partial v}{\partial t} + (v \cdot \nabla) v - \nu \Delta v = u$$

$$v = -\Delta \psi$$

$$\int_D \psi = 0$$

$$v = \left(\frac{\partial \psi}{\partial x_2}, -\frac{\partial \psi}{\partial x_1} \right)$$

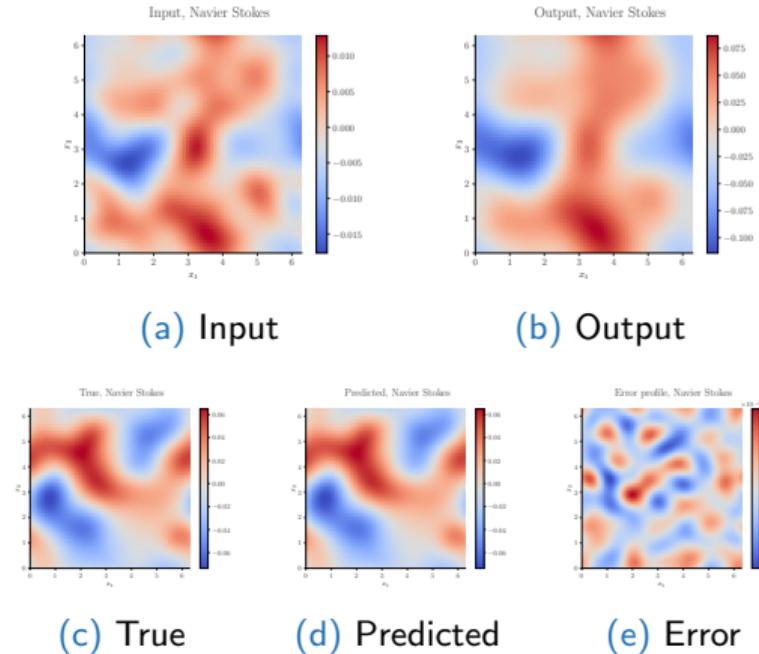
Learn the map from the forcing term f to the vorticity field v at $t = T$:

$$\mathcal{S} : u \mapsto v(\cdot, T)$$

Navier-Stokes

Method	Accuracy
DeepONet	3.63 %
FNO	0.26 %
PCA-Net	2.32 %
Linear method	5.41 %
Kernel method	0.12%

Table: L^2 relative error
on Navier-Stokes.



Accuracy-complexity

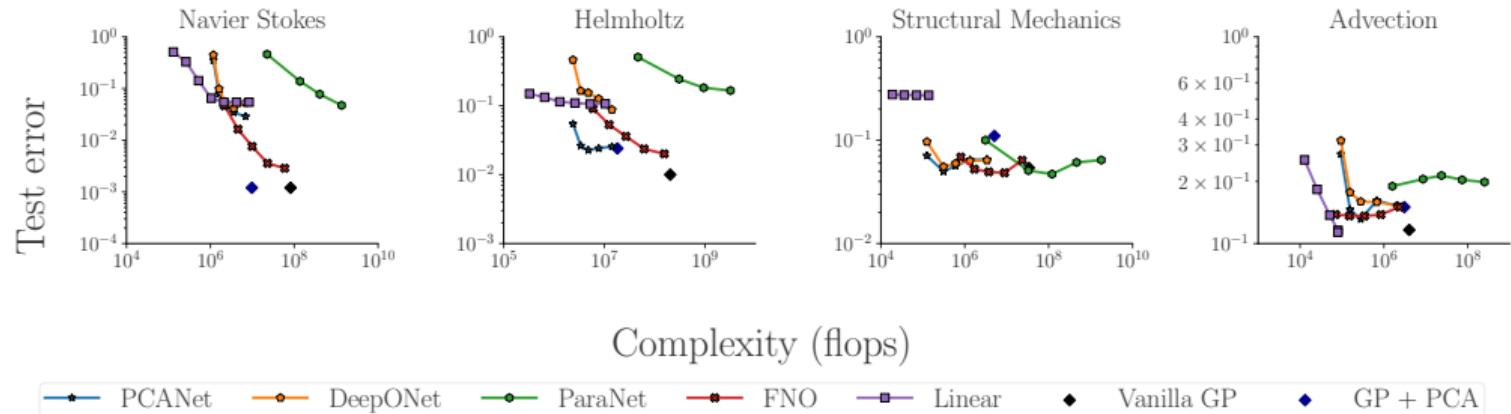


Figure: Cost-accuracy tradeoff at prediction for different operator learning frameworks.

Theory: convergence result

Under a set of technical assumptions, we have the following theorem:

Theorem (Condensed version of Main Theorem)

For all $t' \in (0, t)$,

$$\lim_{n,m \rightarrow \infty} \lim_{N \rightarrow \infty} \sup_{u \in B_R(\mathcal{H}_Q)} \|S(u) - \chi^m \circ \bar{f}_N^{m,n} \circ \phi^n(u)\|_{H^{t'}(D)} \rightarrow 0.$$

Uses the theory of kernel interpolation + reconstruction operators.

Summary of results: accuracy

	Low-data regime			High-data regime			
	Burger's	Darcy problem	Advection I	Advection II	Hemholtz	Structural Mechanics	Navier Stokes
DeepONet	2.15%	2.91%	0.66%	15.24%	5.88%	5.20%	3.63%
POD-DeepONet	1.94%	2.32%	0.04%	n/a	n/a	n/a	n/a
FNO	1.93%	2.41%	0.22%	13.49%	1.86%	4.76%	0.26%
PCA-Net	n/a	n/a	n/a	12.53%	2.13%	4.67%	2.65%
PARA-Net	n/a	n/a	n/a	16.64%	12.54%	4.55%	4.09%
Linear	36.24%	6.74%	$2.15 \times 10^{-13}\%$	11.28%	10.59%	27.11%	5.41%
Best of Matérn/RQ	2.15%	2.75%	$2.75 \times 10^{-3}\%$	11.44%	1.00%	5.18%	0.12%

Table: Summary of numerical results: L^2 relative test error.

We compare the performance against several neural operator frameworks on a range of problems taken from two benchmarking studies ² and find competitive performance.

²[Hoop, Huang, Qian, and Stuart 2022, Lu, Meng, Cai, Mao, Goswami, Zhang, and G. E. Karniadakis 2022]

Our key contributions are:

- A simple, low-cost, and competitive kernel method for operator learning.
- Theoretical guarantees for these methods.
- A general framework for doing operator learning with kernel methods.

Batlle, D, Hosseini, and Owhadi "Kernel methods are competitive for operator learning", *Journal of Computational Physics*, 2024

